

trained

November 13, 2024

```
[1]: import tensorflow as tf
from tensorflow.keras.applications import ResNet50, MobileNetV2
from tensorflow.keras.layers import Input, Dense, Flatten, Dropout, Conv2D,
    ↪MaxPooling2D
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import numpy as np
import warnings
warnings.filterwarnings("ignore", category=UserWarning, module="keras")

# Define ResNet-18 architecture (using ResNet50 with custom layers)
def build_resnet18(input_shape=(32, 32, 3), num_classes=10):
    base_model = tf.keras.applications.ResNet50(
        input_shape=input_shape,
        include_top=False,
        weights=None # No pretrained weights due to smaller input shape
    )
    model = Sequential([
        Input(shape=input_shape),
        base_model,
        Flatten(),
        Dense(256, activation='relu'),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])
    return model

# Define MobileNetV2 architecture
def build_mobilenet(input_shape=(32, 32, 3), num_classes=10):
    base_model = MobileNetV2(input_shape=input_shape, include_top=False,
    ↪weights=None) # No pretrained weights
    model = Sequential([
        Input(shape=input_shape),
        base_model,
```

```

        Flatten(),
        Dense(128, activation='relu'),
        Dropout(0.4),
        Dense(num_classes, activation='softmax')
    ])
    return model

# Define AlexNet architecture
def build_alexnet(input_shape=(32, 32, 3), num_classes=10):
    model = Sequential([
        Input(shape=input_shape),
        Conv2D(96, kernel_size=(3, 3), activation='relu'),
        MaxPooling2D(pool_size=(2, 2)),
        Conv2D(256, kernel_size=(3, 3), activation='relu'),
        MaxPooling2D(pool_size=(2, 2)),
        Flatten(),
        Dense(512, activation='relu'),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])
    return model

```

```

[2]: # Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize pixel values to [0, 1]
y_train = tf.keras.utils.to_categorical(y_train, 10) # One-hot encode labels
y_test = tf.keras.utils.to_categorical(y_test, 10)

# Split training data into training and validation sets
x_train, x_val = x_train[:40000], x_train[40000:]
y_train, y_val = y_train[:40000], y_train[40000:]

```

```

[3]: # Set up data augmentation
datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(x_train)

```

```

[4]: # Function to get a new instance of an optimizer
def get_optimizer(optimizer_name):
    if optimizer_name == 'Adam':
        return Adam(learning_rate=0.0001)
    elif optimizer_name == 'SGD':

```

```

        return SGD(learning_rate=0.001, momentum=0.9)
    else:
        raise ValueError("Optimizer name must be either 'Adam' or 'SGD'")

# Function to train and evaluate each model
def train_and_evaluate(model, model_name, optimizer_name, x_train, y_train,
    ↪x_val, y_val):
    optimizer = get_optimizer(optimizer_name) # Get a new optimizer instance
    model.compile(optimizer=optimizer, loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

    # Use try-except to capture any training errors and ensure cleanup
    try:
        history = model.fit(
            datagen.flow(x_train, y_train, batch_size=128),
            epochs=8,
            validation_data=(x_val, y_val),
            verbose=1
        )
    except Exception as e:
        print(f"Error during training {model_name} with {optimizer_name}: {e}")
        return None, None # Return None to avoid further processing

    plot_history(history, model_name, optimizer_name)
    val_loss, val_accuracy = model.evaluate(x_val, y_val, verbose=0)
    print(f'{model_name} with {optimizer_name}: Validation Accuracy =
    ↪{val_accuracy:.4f}, Validation Loss = {val_loss:.4f}')
    return val_accuracy, val_loss

```

```

[5]: # Function to plot training and validation accuracy and loss
def plot_history(history, model_name, optimizer_name):
    if not history:
        return # Skip if training was not successful

    plt.figure(figsize=(12, 4))

    # Plot accuracy
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{model_name} - {optimizer_name} Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    # Plot loss
    plt.subplot(1, 2, 2)

```

```

plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} - {optimizer_name} Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()

```

```

[6]: # Dictionary to hold model architectures
models = {
    'ResNet-18': build_resnet18(input_shape=(32, 32, 3), num_classes=10),
    'MobileNetV2': build_mobilenet(input_shape=(32, 32, 3), num_classes=10),
    'AlexNet': build_alexnet(input_shape=(32, 32, 3), num_classes=10)
}

# Dictionary to store results
results = {}

# Train each model with each optimizer
for model_name, model in models.items():
    for optimizer_name in ['Adam', 'SGD']: # Specify optimizer names
        print(f'Training {model_name} with {optimizer_name} optimizer...')
        val_accuracy, val_loss = train_and_evaluate(model, model_name,
optimizer_name, x_train, y_train, x_val, y_val)
        if val_accuracy is not None and val_loss is not None:
            results[(model_name, optimizer_name)] = (val_accuracy, val_loss)

```

Training ResNet-18 with Adam optimizer...

Epoch 1/8

313/313 399s 1s/step -

accuracy: 0.1259 - loss: 2.6843 - val\_accuracy: 0.1076 - val\_loss: 2.3954

Epoch 2/8

313/313 392s 1s/step -

accuracy: 0.1956 - loss: 2.1391 - val\_accuracy: 0.2344 - val\_loss: 2.0408

Epoch 3/8

313/313 393s 1s/step -

accuracy: 0.2391 - loss: 2.0421 - val\_accuracy: 0.2970 - val\_loss: 1.9195

Epoch 4/8

313/313 862s 3s/step -

accuracy: 0.2751 - loss: 1.9697 - val\_accuracy: 0.3277 - val\_loss: 1.8377

Epoch 5/8

313/313 391s 1s/step -

accuracy: 0.3036 - loss: 1.9002 - val\_accuracy: 0.3590 - val\_loss: 1.7614


Epoch 6/8

313/313 391s 1s/step -

accuracy: 0.3417 - loss: 1.8170 - val\_accuracy: 0.3737 - val\_loss: 1.7280


Epoch 7/8

313/313                      1248s 4s/step -  
 accuracy: 0.3683 - loss: 1.7584 - val\_accuracy: 0.4013 - val\_loss: 1.6393  
 Epoch 8/8  
 313/313                      390s 1s/step -  
 accuracy: 0.3960 - loss: 1.6875 - val\_accuracy: 0.4353 - val\_loss: 1.5698



ResNet-18 with Adam: Validation Accuracy = 0.4353, Validation Loss = 1.5698  
 Training ResNet-18 with SGD optimizer...

Epoch 1/8  
 313/313                      394s 1s/step -  
 accuracy: 0.3543 - loss: 1.7890 - val\_accuracy: 0.3235 - val\_loss: 2.0626  
 Epoch 2/8  
 313/313                      389s 1s/step -  
 accuracy: 0.4168 - loss: 1.6101 - val\_accuracy: 0.4053 - val\_loss: 1.6533  
 Epoch 3/8  
 313/313                      389s 1s/step -  
 accuracy: 0.4390 - loss: 1.5449 - val\_accuracy: 0.4732 - val\_loss: 1.4585  
 Epoch 4/8  
 313/313                      388s 1s/step -  
 accuracy: 0.4668 - loss: 1.4886 - val\_accuracy: 0.4285 - val\_loss: 1.5386  
 Epoch 5/8  
 313/313                      389s 1s/step -  
 accuracy: 0.4836 - loss: 1.4412 - val\_accuracy: 0.4698 - val\_loss: 1.5027  
 Epoch 6/8  
 313/313                      6683s 21s/step -  
 accuracy: 0.4980 - loss: 1.4017 - val\_accuracy: 0.5037 - val\_loss: 1.4115  
 Epoch 7/8  
 313/313                      17073s 55s/step -  
 accuracy: 0.5200 - loss: 1.3580 - val\_accuracy: 0.5179 - val\_loss: 1.3328  
 Epoch 8/8  
 313/313                      6366s 20s/step -  
 accuracy: 0.5316 - loss: 1.3164 - val\_accuracy: 0.5455 - val\_loss: 1.2640



ResNet-18 with SGD: Validation Accuracy = 0.5455, Validation Loss = 1.2640

Training MobileNetV2 with Adam optimizer...

Epoch 1/8

313/313 34s 89ms/step -

accuracy: 0.1184 - loss: 2.4307 - val\_accuracy: 0.0997 - val\_loss: 2.3028

Epoch 2/8

313/313 29s 91ms/step -

accuracy: 0.1651 - loss: 2.2193 - val\_accuracy: 0.1016 - val\_loss: 2.3031

Epoch 3/8

313/313 28s 89ms/step -

accuracy: 0.1993 - loss: 2.1410 - val\_accuracy: 0.0977 - val\_loss: 2.3030

Epoch 4/8

313/313 28s 89ms/step -

accuracy: 0.2284 - loss: 2.0500 - val\_accuracy: 0.0977 - val\_loss: 2.3049

Epoch 5/8

313/313 28s 89ms/step -

accuracy: 0.2500 - loss: 1.9932 - val\_accuracy: 0.0977 - val\_loss: 2.3073

Epoch 6/8

313/313 27s 88ms/step -

accuracy: 0.2796 - loss: 1.9192 - val\_accuracy: 0.0977 - val\_loss: 2.3095

Epoch 7/8


313/313 27s 86ms/step -

accuracy: 0.3029 - loss: 1.8687 - val\_accuracy: 0.0977 - val\_loss: 2.3122

Epoch 8/8

313/313 27s 86ms/step -

accuracy: 0.3169 - loss: 1.8245 - val\_accuracy: 0.0977 - val\_loss: 2.3130



MobileNetV2 with Adam: Validation Accuracy = 0.0977, Validation Loss = 2.3130

Training MobileNetV2 with SGD optimizer...

Epoch 1/8

313/313 30s 83ms/step -

accuracy: 0.3335 - loss: 1.8000 - val\_accuracy: 0.0977 - val\_loss: 2.3131

Epoch 2/8

313/313 26s 84ms/step -

accuracy: 0.3776 - loss: 1.6747 - val\_accuracy: 0.1016 - val\_loss: 2.3152

Epoch 3/8

313/313 27s 86ms/step -

accuracy: 0.4191 - loss: 1.5798 - val\_accuracy: 0.1016 - val\_loss: 2.3202

Epoch 4/8

313/313 27s 85ms/step -

accuracy: 0.4362 - loss: 1.5464 - val\_accuracy: 0.0952 - val\_loss: 2.3256

Epoch 5/8

313/313 27s 86ms/step -

accuracy: 0.4633 - loss: 1.4713 - val\_accuracy: 0.0952 - val\_loss: 2.3176

Epoch 6/8

313/313 27s 87ms/step -

accuracy: 0.4779 - loss: 1.4452 - val\_accuracy: 0.0952 - val\_loss: 2.3273

Epoch 7/8


313/313 27s 86ms/step -

accuracy: 0.4992 - loss: 1.3990 - val\_accuracy: 0.0960 - val\_loss: 2.3371

Epoch 8/8

313/313 27s 87ms/step -

accuracy: 0.5061 - loss: 1.3814 - val\_accuracy: 0.1242 - val\_loss: 2.3324



MobileNetV2 with SGD: Validation Accuracy = 0.1242, Validation Loss = 2.3324

Training AlexNet with Adam optimizer...

Epoch 1/8

313/313 27s 84ms/step -

accuracy: 0.2394 - loss: 2.0681 - val\_accuracy: 0.3915 - val\_loss: 1.6855

Epoch 2/8

313/313 27s 85ms/step -

accuracy: 0.4034 - loss: 1.6423 - val\_accuracy: 0.4592 - val\_loss: 1.5150

Epoch 3/8

313/313 27s 85ms/step -

accuracy: 0.4521 - loss: 1.5130 - val\_accuracy: 0.5023 - val\_loss: 1.3888

Epoch 4/8

313/313 26s 84ms/step -

accuracy: 0.4798 - loss: 1.4487 - val\_accuracy: 0.5231 - val\_loss: 1.3531

Epoch 5/8

313/313 26s 83ms/step -

accuracy: 0.4998 - loss: 1.3921 - val\_accuracy: 0.5163 - val\_loss: 1.3640

Epoch 6/8

313/313 26s 83ms/step -

accuracy: 0.5165 - loss: 1.3449 - val\_accuracy: 0.5571 - val\_loss: 1.2572

Epoch 7/8

313/313 26s 83ms/step -


accuracy: 0.5368 - loss: 1.3021 - val\_accuracy: 0.5753 - val\_loss: 1.2145

Epoch 8/8

313/313 26s 83ms/step -

accuracy: 0.5425 - loss: 1.2810 - val\_accuracy: 0.5982 - val\_loss: 1.1534





AlexNet with Adam: Validation Accuracy = 0.5982, Validation Loss = 1.1534

Training AlexNet with SGD optimizer...

Epoch 1/8

313/313 25s 80ms/step -

accuracy: 0.5293 - loss: 1.3198 - val\_accuracy: 0.5751 - val\_loss: 1.2186

Epoch 2/8

313/313 25s 80ms/step -

accuracy: 0.5474 - loss: 1.2747 - val\_accuracy: 0.5879 - val\_loss: 1.1595

Epoch 3/8

313/313 25s 80ms/step -

accuracy: 0.5612 - loss: 1.2352 - val\_accuracy: 0.6069 - val\_loss: 1.1196

Epoch 4/8

313/313 25s 81ms/step -

accuracy: 0.5664 - loss: 1.2376 - val\_accuracy: 0.6004 - val\_loss: 1.1360

Epoch 5/8

313/313 25s 81ms/step -

accuracy: 0.5676 - loss: 1.2242 - val\_accuracy: 0.6119 - val\_loss: 1.1052

Epoch 6/8

313/313 25s 81ms/step -

accuracy: 0.5697 - loss: 1.2130 - val\_accuracy: 0.6003 - val\_loss: 1.1494

Epoch 7/8


313/313 26s 82ms/step -

accuracy: 0.5747 - loss: 1.1983 - val\_accuracy: 0.6223 - val\_loss: 1.0896

Epoch 8/8

313/313 25s 81ms/step -

accuracy: 0.5800 - loss: 1.1951 - val\_accuracy: 0.6129 - val\_loss: 1.1107



AlexNet with SGD: Validation Accuracy = 0.6129, Validation Loss = 1.1107

```
[7]: # Find the best model and optimizer based on validation accuracy
if results:
    best_model_info = max(results, key=lambda x: results[x][0])
    best_model_name, best_optimizer_name = best_model_info
    best_model = models[best_model_name]
    best_optimizer = get_optimizer(best_optimizer_name)

    print(f"\nBest Model: {best_model_name} with {best_optimizer_name}
    ↪optimizer")

    # Recompile best model with best optimizer and evaluate on the test set
    best_model.compile(optimizer=best_optimizer,
    ↪loss='categorical_crossentropy', metrics=['accuracy'])
    test_loss, test_accuracy = best_model.evaluate(x_test, y_test, verbose=0)
    print(f'\nTest Accuracy of Best Model ({best_model_name} with
    ↪{best_optimizer_name}): {test_accuracy:.4f}')
else:
    print("No valid results to display.")
```

Best Model: AlexNet with SGD optimizer

Test Accuracy of Best Model (AlexNet with SGD): 0.6115

```
[ ]: 
```

```
[ ]: ]
```

```
[ ]: 
```